Loss-Dynamics-Aware Multi-Agent Cooperative Active Sampling System for Efficient Robotic Online Learning

Anonymous Author(s)
Submission Id: 320*

Abstract

Online training tasks rely on training input consecutively sampled from the real world to refine their training model, and automated mobile devices (robots) shouldering the responsibility of sampling has the potential to further enhance their training performance by actively sampling in reaction to the real-time training quality, besides saving human labor. However, to achieve such active sampling faces two major gaps. First, as the training quality is typically computation-intensive to get for a robot (e.g., validating), where to move for the robot to boost training performance? Second, since the consecutive samples from nearby states (e.g., position, orientation) of a robot are too similar and would limit the training information gain, how to move for better efficiency?

We observe that real-time training loss is spatially and temporally related to locations in the environment and implies the level of information gain on further sampling of these locations. The second gap can be overcome by cooperative multi-view sampling from multiple robots, since they are naturally distant in state space. Based on these observation, we choose to build an online training application named MIAS (Multi-robot Implicit Active SLAM) that drives multiple robots to actively and cooperatively sample the environment in real-time reaction to the quality of the training implicit SLAM model. Evaluation shows that MIAS with three robots speeds up the implicit SLAM tasks not only by up to xxX compared to the baselines with three robots, but also by xxX compared to MIAS with single robot.

CCS Concepts: • **Computer systems organization** → **Robotic control**.

Keywords: Online Training, Active Learning, Multi-Agent System, 3D Reconstruction

1 Introduction

Online training tasks (e.g., implicit SLAM, domain adaptation, long term learning) take unlabelled training input consecutive sampled from the real world to refine their training model for better performance in the changing real world environments. The sampling of such training input typically relies on human labor (e.g., the RGBD image sequences captured by hand-held camera in implicit SLAM); offloading such online training tasks to automated mobile devices (robots equipped with GPUs) not only enables automated sampling, which saves human labor, but also has the potential to enable active sampling in reaction to the real-time training quality to boost training performance (e.g., sample the training input with highest training information gain).

However, enabling active sampling for online training tasks deployed over robots faces two major gaps. First, where to move for higher training information gain? Typically, quantifying the training quality of the training model relies on statistical methods such as validating, which is computation-intensive for robots and infects the ongoing online training task. Without the real-time statistics of training quality, it is difficult to decide the destination for the robot to mobile in active sampling.

Second, how to move for higher training information gain? The consecutive sampling of a robot suffers the problem of sampling locality: while the robot is moving in local state space (i.e., nearby position and orientation), the consecutive samples are often too similar, limiting training information gain. We visualize in Fig.**??** the training information gain of consecutive training input about an area of interest (e.g., a wall that has high potential training information gain) after a robot sampling this area once in an implicit SLAM task and we can find that there is a lowland of training information gain around the starting state of the robot. As a result, to complete sampling of an area of interest, the robot has to move in and out the lowland of training information gain during sampling, lowering the average training information gain per sample.

To overcome to above gaps, we observe that real-time training loss is spatially and temporally related to locations in the environment and implies the level of information gain on further sampling of these locations. ALthough training loss does not directly reveal the statistics of training model quality such as accuracy, but it is semantically related with state space when the robot sampled the corresponding training input. A higher training loss compared with others reveals the training model is more inaccurate against the corresponding training input and higher possible level of training information gain of further sampling in the corresponding state space, which has the potential to guide the navigation for active sampling.

For the second gap, we find that multiple robots are naturally distant in state space and their sampling from multiple perspectives (multi-view sampling) has the potential the skip the lowland of training information and mitigate the locality problem. With multi-view sampling, it is possible that after one robot samples an area of interest, other distant robots would shoulder the responsibility of further sampling the same area (if loss is still high) beyond the lowland of training information gain and the first robot could quickly switch its target and avoid the lowland.

With the above ideas, we take the first step to build a multi-robot active online training application and choose implicit SLAM as the main workload, naming MIAS, Multi-view Implicit Active SLAM. The implicit SLAM builds dense mesh of the surrounding environment in real time by optimizing both the localization of the state (position and orientation) of the robot (i.e., tracking) and an implicit representation of the dense mesh (i.e., mapping) over consecutive RGBD images sampled from robot cameras. Traditional methods to automate the sampling for SLAM are unfit for implicit SLAM: first, they trace the explicit representation of the map (e.g., the dense mesh) to estimate the quality of tracking and mapping for decision making, which could only be obtained by computation intensive validation of the implicit representation in implicit SLAM, slowing down the decision making; second, they typically avoid multi-view sampling of multiple robots so that the robots would sample different areas to optimize coverage, suffering the sampling locality problem.

But with MIAS, we can achieve decision making in real-time reaction to online training quality by tracing the change of loss level: when the training loss for tracking increases, the state of the robot is inaccurate and we control the robot to sample the previously sampled area for re-localization; when the training loss for mapping increases, we mark the corresponding areas as places of interest for further sampling; when both losses are low, we explore new areas. The sampling of places of interest is further accelerated by multi-view sampling beyond the lowland of information gain. As a result, both training quality and training information gain per sample are optimized.

We implemented MIAS on four-wheel robots equipped with RGBD cameras and state-of-the-art (SOTA) mobile GPU chips. The implicit representation of dense mesh of implicit SLAM is distributedly trained among the robots over SOTA distributed training library optimized for mobile devices with gradient compression. We find that the cost for distributed training (e.g., time for compressing and communicating the gradients) is little for a single robot, accounting for only xx% time of the computation of gradients. We compare MIAS with two traditional active SLAM methods and evaluated over a team of three robots in both habitat-simulated environments and real-world environments. Evaluation shows that:

- MIAS accelerates the active implicit SLAM process. With the same time cost, MIAS increased the accuracy of the built dense mesh by xxX to xxX compared with the baselines. When reaching the same high accuracy of the built dense mesh, MIAS reduced the total time cost by xx% to xx%.
- N robots involved brings N+ times acceleration. When reaching the same high accuracy of the built dense mesh, MIAS with two / three robots reduced the total time cost by xx% to xx% / xx% to xx% compared with the baselines and MIAS with one robot. With the same time cost, MIAS with two / three robots increased the accuracy of the built dense mesh by xxX to xxX / xxX to xxX.

Our major contributions are the paradigm of enabling active sampling for online training tasks deployed over robots and the real-world multi-robot active implicit SLAM application, MIAS. The paradigm would benefit various online training tasks deployed on robots such as long term training or domain adaptation in the filed by bridging the real-time interaction between online training and robot mobility, so that the mobility of robots can better serve for optimizing the training quality of online training tasks via action decision making based on real-time training loss and optimizing the training information gain per sample via multi-view sampling. The application MIAS not only automatically samples the environment to build dense mesh of the environment in real time that saves human labor, but also optimizes the quality of dense mesh and the training information gain per sample by active sampling, leading to acceleration of dense mesh building beyond the number of robots involved.

The rest of this paper is organized as follows:...

2 Background & RELATED WORK

2.1 Online Training

2.2 Active Learning

2.3 Active Sampling in Traditional Robotics

3 Overview

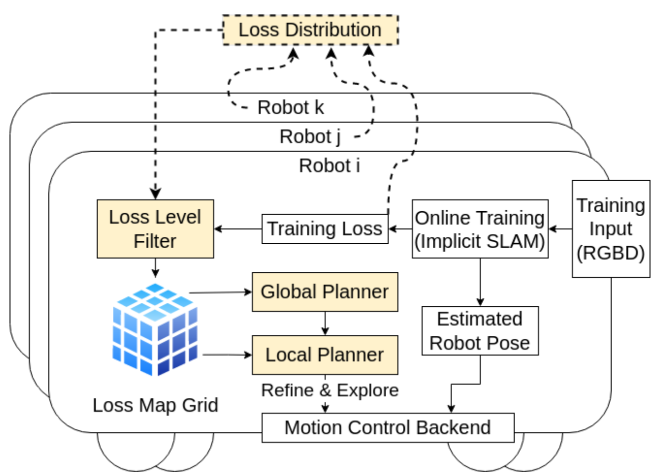

Figure 1. Overview of MIAS

This chapter presents the architecture of MIAS and gives an overview of how MIAS achieves real-decision making in reaction to real-time training quality and cooperative multi-view sampling across robots.

4 Evaluation

baselines: badge, uncertainty, random

4.1 End-to-End Performance

4.1.1 Different Number of Robots. figures: y accuracy; x time; 1 robot; 2 robots; 3 robots

fact

ours all better than the baselines; more robot, diff larger badge, uncertainty > random, but due to ... with our system, ...

4.1.2 Different Size of Scene. figure: y accuracy; x time; small scene; medium scene; large scene

facts

ours advantages remain, larger scene, better badge, uncertainty > random, but due to ...

our can be better because...

4.2 Ablation Study

References

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009